

# RADICLE REGISTRY

## SPECIFICATION

VERSION 0.1

MONADIC<sup>†</sup>

### 1. TRANSACTIONS

All transactions on the registry take the form  $\text{transaction}(arg_1, \dots, arg_n)_\sigma$ , where  $arg_1, \dots, arg_n$  are the *inputs* and  $\sigma$  is the EdDSA signature of the author of the transaction. Transactions always have an *author* and an *origin* (formally  $\alpha$ ), which is the author's account.

Transactions can be uniquely identified by their *hash*. The set of all *known* transactions is  $\mathcal{T}$  (the “ledger”), and the set of all known transaction hashes is  $\mathcal{T}_{\text{hash}}$ .

### 2. ACCOUNTS

An account  $A$  is a tuple:

$$A = \langle A_{\text{id}}, A_{\text{nonce}}, A_{\text{bal}} \rangle$$

DEFINITION

- $A_{\text{id}}$  is the unique account identifier obtained by hashing the account owner's public key,
- $A_{\text{nonce}}$  is a number which starts at 0 and is incremented every time a transaction originates from this account.
- $A_{\text{bal}}$  is the account's balance in the smallest denomination, and  $A_{\text{bal}} \in \mathbb{N}_{\geq 0}$ .

The set of all accounts is  $\mathcal{A}$ . Accounts are never created or destroyed, rather, if they have never been used to transact, they have an initial state of:

$$A = \langle A_{\text{id}}, 0, 0 \rangle$$

Hence, for all valid account ids, there exists an account with that id. In other words,  $\forall a \in A_{\text{id}} (A \in \mathcal{A})$ .

Note that accounts can *never be removed*, since that would violate the invariant that nonces are only ever incremented, and removing an account is equivalent to setting  $A_{\text{nonce}}$  and  $A_{\text{bal}}$  to 0.

**2.1. Root accounts.** Some accounts are considered *privileged*. These ‘root’ accounts, formally  $\mathcal{A}_R \subset \mathcal{A}$  are authorized to conduct certain transactions that are only valid when originating from these accounts.

The set of accounts in  $\mathcal{A}_R$  is defined at *genesis*, and may not be further modified in the initial protocol.

**2.2. Transferring value.** The act of transferring coins between two accounts:

$$\text{transfer}(A_{\text{id}}, v)_\sigma$$

which will transfer value from the transaction origin  $\alpha$  to account  $A$ .

INPUTS

- $A_{\text{id}}$  is the account id of the *receiver* of the transfer,
- $v$  is the value or ‘balance’ to transfer from the origin to the receiver, in the smallest denomination.

VALIDATION

- The transfer balance is positive, or  $v \geq 1$ ,
- The origin's balance minus any transaction fee is  $\geq v$ .

OUTPUTS

- $v$  is debited from the origin and credited to  $A$ .

### 3. PROJECTS

A project  $P$  is a tuple:

$$P = \langle P_{\text{id}}, P_{\text{k}}, P_{\text{account}}, P_{\text{contract}}, P_{\text{proof}}, P_{\text{meta}} \rangle$$

DEFINITION

- $P_{\text{id}}$  is the unique project identifier, defined as  $\langle P_{\text{name}}, P_{\text{domain}} \rangle$ ,
- $P_{\text{k}}$  is the current project *checkpoint* (See §3.4),
- $P_{\text{account}}$  is the project account or *fund*,
- $P_{\text{contract}}$  is the project contract, which governs permissions around the project, as well as its fund.

It can be described as a function:

$$f : \mathcal{T} \rightarrow \mathcal{O}$$

where  $\mathcal{T}$  is any known transaction, and  $\mathcal{O}$  is the function's codomain described by:

$$\{\top, \perp\} \cup [\langle A_{\text{id}}, \mathbb{N} \rangle] \cup \mathbb{N}.$$

An output in  $\{\top, \perp\}$  is reserved for *boolean* contracts, where  $\top$  signifies a transaction  $t \in \mathcal{T}$  is *authorized* to execute by the contract, and  $\perp$  means it is *unauthorized*. Note that a transaction can be verified and included in the transaction ledger  $\mathcal{T}$  yet still be unauthorized to run.

An output in  $[\langle A_{\text{id}}, \mathbb{N} \rangle]$  is reserved for transactions that should trigger a distribution of funds to one or more contract-specified beneficiaries, where

<sup>†</sup>Monadic, team@monadic.xyz.

$A_{id}$  is the account of one of the beneficiaries, and  $\mathbb{N}$  is the value to transfer.

An output in  $\mathbb{N}$  is reserved for transactions that should trigger a transfer of funds to an *a priori known* beneficiary that shouldn't be determined by the contract. For example, the transaction origin.

- $P_{proof}$  is the proof that was supplied during proposal, verifying the owner's authority over the project,
- $P_{meta}$  is a dictionary of additional metadata to associate with the project, for example the RADICLE *project id*.  $P_{meta}$  is *immutable* once defined.

Projects are proposed with the **propose-project** transaction and subsequently registered with the **accept-project** transaction. The set of all registered projects is  $\mathcal{P}$ . The set of all proposals is  $\rho$ .

**3.1. Project proposal.** The act of proposing a project for registration under a unique name and domain:

**propose-project**( $P_{id}, P_k, P_{contract}, P_{proof}, P_{meta}$ ) $_{\sigma}$

This transaction requires a deposit  $\mathcal{D}_{propose-project}$ .

INPUTS

- $P_{id} = \langle P_{name}, P_{domain} \rangle$
- $P_{name}$  is the unique name being requested, where
- $P_{domain}$  is the domain under which  $P_{name}$  is being registered, which together form the unique identifier  $P_{id}$ ,
- $P_k$  is the id of the initial *checkpoint* associated with this project, formally  $k_0$ . This checkpoint must always remain in the project ancestry,
- $P_{contract}$  is the initial project contract that includes the initial permission set around the project,
- $P_{proof}$  is a byte array of up to 4096 bytes supplied to prove the legitimacy of this project proposal—it is verified *off-registry* during project approval (See §3.2),
- $P_{meta}$  is a dictionary of metadata to be associated with  $P$ . For example, identities on other platforms. Note that once submitted, the metadata is immutable.

VALIDATION

- $P_{name}$  must be unique, i.e. not currently registered under  $P_{domain}$ , between 1 and 32 characters long, and valid UTF-8—note that it is perfectly valid for multiple proposals to co-exist with the same  $P_{id}$ .
- $P_{domain}$  must be an existing domain,
- $P_{proof}$ 's size is  $\leq 4096$  bytes,
- $P_k$  must represent an existing checkpoint,
- $\alpha_{bal} \geq \mathcal{D}_{propose-project}$ .

OUTPUTS

- $P \in \rho$ , where  $\rho$  is the set of proposals that are in a “pending” state, waiting to be accepted or rejected (See §3.2). Note that  $\rho$  and  $\mathcal{P}$  are disjoint,
- $\alpha_{bal'} = \alpha_{bal} - \mathcal{D}_{propose-project}$ .

**3.2. Accepting and rejecting a project.** The act of accepting or rejecting a project being registered:

**accept-project**( $t_{hash}$ ) $_{\sigma}$

or

**reject-project**( $t_{hash}$ ) $_{\sigma}$

INPUTS

- $t_{hash}$  is the *transaction hash* of the **propose-project** transaction  $t$  of a project  $P$  being accepted or rejected.

VALIDATION

- The transaction *origin* is a member of  $\mathcal{A}_R$ ,
- $t_{hash}$  must be the hash of an existing transaction of type **propose-project**. In other words,  $P \in \rho$ , where  $P$  is the project being registered.

For **accept-project**,

OUTPUTS

- $P \in \mathcal{P}$
- $P \notin \rho$
- $P_{account} = \langle A_{id}, 0, 0 \rangle$

For **reject-project**,

OUTPUTS

- $P \notin \mathcal{P}$
- $P \notin \rho$

**3.3. Project proposal withdrawal.** It's possible to withdraw a project from proposal if it hasn't been accepted or rejected yet with:

**withdraw-project**( $t_{hash}$ ) $_{\sigma}$ .

This in turn returns the deposit made on **propose-project**.

INPUTS

- $t_{hash}$  is the transaction hash of the **propose-project** transaction  $t$  for some project  $P$  that should be withdrawn.

VALIDATION

- $P \in \rho$

OUTPUTS

- $P \notin \rho$
- $\alpha_{bal'} = \alpha_{bal} + \mathcal{D}_{propose-project}$ .

**3.4. Checkpointing.** The act of notarizing a project's state and updating the network graph:

**checkpoint**( $K_{parent}, K_{hash}, K_{version}, K_{contribs}, K_{deps}$ ) $_{\sigma}$

Checkpoints within the scope of a single project form a chain going from the latest, or “current” checkpoint  $k_{n-1}$  to the first and original checkpoint  $k_0$ . Checkpoints are identified by their transaction hash, so  $k \in T_{hash}$ .

From the perspective of  $k_0$ , we can talk of a checkpoint *tree*, since due to their nature, they are able to represent branching. Hence, the original checkpoint  $k_0$  is also called the *root* checkpoint.

INPUTS

- $K_{\text{parent}}$  is the *id* of the previous or ‘parent’ checkpoint,
- $K_{\text{hash}}$  is the new hash of the project state,
- $K_{\text{version}}$  is the current version of the project,
- $K_{\text{contribs}}$  is the list of contributions since  $K_{\text{parent}}$ ,
- $K_{\text{deps}}$  is the list of dependency updates since the  $K_{\text{parent}}$ .

## VALIDATION

- $K_{\text{parent}}$  refers to an existing checkpoint in the registry, or is  $\emptyset$ .
- $K_{\text{hash}}$  is a valid hash that hasn’t been used in a parent checkpoint.
- $K_{\text{version}}$  is a string between 1 and 32 bytes long that *may* have been used in a previous project checkpoint.
- $K_{\text{contribs}}$  is a valid contribution list (See §3.4.1).
- $K_{\text{deps}}$  is a valid dependency update list (See §3.4.2).

3.4.1. *Contributions.* The list  $K_{\text{contribs}}$  supplied to the checkpoint transaction is of the form:

$$K_{\text{contribs}} = [\langle C_{\text{parent}}, C_{\text{hash}}, C_{\text{author}}, C_{\text{sig}} \rangle],$$

## DEFINITION

- $C_{\text{parent}}$  is the hash of the parent contribution, or  $\emptyset$  if this is the first contribution of the first checkpoint of the project.
- $C_{\text{hash}}$  is the hash of the corresponding commit,
- $C_{\text{author}}$  is the public signing key of the commit referred to by  $C$ ,
- $C_{\text{sig}}$  is the author’s GPG signature.

## VALIDATION

- $C_{\text{parent}}$  is a valid SHA-1 hash or  $\emptyset$  if this is the first contribution. Note that if  $C$  is  $K_{\text{contribs}}$ ’s first item, and  $C'$  is the *last* item of the *parent* checkpoint’s contributions list, then  $C'_{\text{hash}}$  and  $C_{\text{parent}}$  must be equal, such that no gaps between contributions exist.
- $C_{\text{hash}}$  is a valid SHA-1 hash,
- $C_{\text{author}}$  is the creator of  $C_{\text{sig}}$ ,
- $C_{\text{sig}}$  is a valid signature of  $C_{\text{hash}}$ .

Because all changes to a project’s source code are described in checkpoints, it is possible to reconstruct a full hash-linked list of contributions for the entire project. When cross-referenced with the project’s repository, this constitutes a complete historical record of who authored what code. This ensures the project history is auditable and tamper-proof, while providing fundamental information to for the network graph  $\mathcal{N}$ . Note that only contribution *metadata* is stored on-chain.

3.4.2. *Dependency updates.* Conceptually, a project  $P$  depends on another project  $P'$  if it is an “input” to  $P$  in some way:  $P$  references  $P'$  or parts of  $P'$  in its source code, or  $P'$  is a build/test dependency.

The dependency update list  $P_{\text{deps}}$  is a list of *dependency updates*, one of:

$$\begin{cases} \text{depend}(P'_{\text{id}}, P'_{\text{version}}) \\ \text{undepend}(P'_{\text{id}}, P'_{\text{version}}) \end{cases}$$

which refer to the project  $P'$  at a specific version  $P'_{\text{version}}$ . The **depend** update adds a new dependency while the **undepend** update removes a dependency. The updates are processed in order with **depend** only being valid if it adds a dependency that the project does not already have and **undepend** only being valid for current dependencies. The checkpoint is invalid if the update list contains duplicates.

## VALIDATION

- $P'_{\text{id}}$  must be a valid project id, but *does not* have to refer to an existing id in the registry. This allows dependent projects to checkpoint dependencies that have not yet been registered.
- $P'_{\text{version}}$  must be a valid version string, but *does not* have to refer to an existing version of  $P'$ . This allows dependent projects to checkpoint before their dependencies.

As a project maintainer, adding a dependency signals a variety of things depending of the nature of the project:

- They have verified that  $P$  indeed depends on this specific version of  $P'$ .
- That  $P'$  is suitable as a dependency for  $P$ , e.g. if  $P$  has very high security requirements, that  $P'$  fulfills these.

Since contributions to a project carry additional weight—potentially increasing a project’s rank—there is an incentive for maintainers to checkpoint their projects regularly. Similarly, adding dependencies may increase connectivity in the network graph, which may in turn indirectly improve a project’s rank.

3.5. **Setting the project checkpoint.** The act of updating the project to point to a new checkpoint:

$$\text{set-checkpoint}(P_{\text{id}}, k')_{\sigma}$$

which updates  $P_{\text{checkpoint}}$  from  $k$  to  $k'$ .

## INPUTS

- $P_{\text{id}}$  is the id of the project being updated,
- $k'$  is the id of the checkpoint the project should be associated to.

## VALIDATION

- $P_{\text{id}}$  refers to a project that has been accepted in the registry,
- $k'$  is a checkpoint which has the original project checkpoint  $k_0$  in its ancestry.
- $P_{\text{contract}}(\text{set-checkpoint}(P_{\text{id}}, k')) \equiv \top$

## OUTPUTS

- $P_k = k'$

Note that the semantics of this transaction allows for projects to revert to a previous checkpoint, or to adopt a “fork”, as long as the new checkpoint shares part of its ancestry with the previous checkpoint.

**3.6. The project fund.** Each project has an associated account  $P_{\text{account}}$  called the *fund*. To use that account to fund maintenance or other projects, the fund transaction is used:

$$\text{fund}(P_{\text{id}}, A_{\text{id}}, v)_{\sigma}$$

#### INPUTS

- $P_{\text{id}}$  is the id of the project that should initiate the value transfer.
- $A_{\text{id}}$  is the id of the account that should receive the value.
- $v$  is the value being transferred.

#### VALIDATION

- $p_{\text{bal}} \geq v$ , where  $p = P_{\text{account}}$ ,
- $P_{\text{contract}}(\text{fund}(P_{\text{id}}, A_{\text{id}}, v)) \equiv \top$ .

#### OUTPUTS

- $A_{\text{bal}'} = A_{\text{bal}} + v$

**3.7. The project contract.** Every project  $P$  in the registry has a contract denoted  $P_{\text{contract}}$ . The way this contract is invoked is through transactions that act on  $P$ . For example, the fund transaction (§??) which transfers value out of a project is always validated by the project contract before it is authorized to execute.

A contract is made of a set of *rules* that each handle a specific action relating to the project. In the fund example, the fund *rule* would be invoked to determine the outcome of the transaction.

**3.7.1. Updating the project contract.** The act of updating a project’s contract:

$$\text{set-contract}(P_{\text{id}}, c)_{\sigma}$$

#### INPUTS

- $P_{\text{id}}$  is the id of the project,
- $c$  is the new contract.

#### VALIDATION

- $P \in \mathcal{P}$ ,
- $P_{\text{contract}}(\text{set-contract}(P_{\text{id}}, c)) \equiv \top$

#### OUTPUTS

- $P_{\text{contract}} = c$

## 4. NAMES

**4.1. Registering a domain.** The act of registering a top-level domain:

$$\text{register-domain}(\text{domain})_{\sigma}$$

#### INPUTS

- *domain* is the unique domain being registered.

#### VALIDATION

- The transaction origin  $\alpha$  is a member of  $\mathcal{A}_R$ ,
- *domain* must be available for registration, between 1 and 32 characters long, and valid UTF-8.

For example,

$$\text{register-domain}(\text{crates})_{\sigma}$$

## 5. IDENTITY

**5.1. Identifying as a contributor.** The act of identifying yourself as a contributor, by linking a public key used to sign project contributions, to an account in the registry:

$$\text{identify}(I_{\text{pk}}, I_{\text{proof}})_{\sigma}$$

#### INPUTS

- $I_{\text{pk}}$  is the public key that is to be associated with the *origin* account  $\alpha$  if this transaction succeeds.
- $I_{\text{proof}}$  is a proof verifying that the transaction author owns  $I_{\text{pk}}$ .

#### VALIDATION

- $I_{\text{pk}}$  is not already associated with an account,
- $I_{\text{proof}}$  is  $\alpha_{\text{id}}$  signed by the secret key  $sk$  that  $I_{\text{pk}}$  was derived from. In other words,

$$I_{\text{proof}} = \text{encrypt}(\text{hash}(\alpha_{\text{id}}), sk)$$

which is valid if

$$\text{decrypt}(I_{\text{proof}}, I_{\text{pk}}) \equiv \text{hash}(\alpha_{\text{id}})$$

**5.2. Forgetting identities.** When a public key associated with the identify transaction is lost or no longer used, the following transaction will ‘forget’ the association:

$$\text{forget}(I_{\text{pk}})_{\sigma}$$

## APPENDIX

**Alternative contribution model.** It’s possible to drastically reduce the size of checkpoints by merkleizing the commit history and making contribution claiming an explicit action on the part of the contributor.

First, we change the checkpoint transaction’s  $K_{\text{contribs}}$  argument from a *list* of commits to a single *hash*. This hash is the Merkle root of the commit list that was  $K_{\text{contribs}}$ , with  $\langle C_{\text{hash}}, C_{\text{author}}, C_{\text{sig}} \rangle$  at each leaf:

$$K_{\text{contribs}} = \text{hash}([\langle C_{\text{hash}}, C_{\text{author}}, C_{\text{sig}} \rangle]),$$

where *hash* is a Merkle hash function.  $K_{\text{contribs}}$  is now constant-sized. Note that in this alternative model,  $C_{\text{parent}}$  is no longer used, since it can’t be verified.

Then, we ask contributors to claim contributions if they wish to receive rewards, with a new claim-contribution transaction:

$$\text{claim-contribution}(P_{\text{id}}, C_1 \dots C_n, C_{\text{proof}})_{\sigma}$$

#### INPUTS

- $P_{\text{id}}$  is the id of the project  $P$  to claim the contributions on.

- $C_1 \dots C_n$  is the list of commits to claim, where  $C = \langle C_{\text{hash}}, C_{\text{sig}} \rangle$ , and  $C_{\text{author}} = \alpha$ , the author of the transaction. Note that  $C_{\text{author}}$  is not actually included in the transaction, since it is assumed to always be  $\alpha$ .
- $C_{\text{proof}}$  is the Merkle proof asserting that  $C_1 \dots C_n$  are included in the  $K_{\text{contribs}}$  set. It is expected that  $C_{\text{proof}}$  can use structural sharing to compress the Merkle paths such that the transaction does not grow linearly in  $n$ . Formally, each  $C_{\text{proof}}$  is an ordered set  $\{\pi_1 \dots \pi_n\}$  of Merkle paths associated with  $C_1 \dots C_n$ .

## VALIDATION

- $C_{\text{proof}}$  is valid if each path  $\pi$  hashes to the same root hash  $r$ , and there is a checkpoint  $K$  in  $P_k$ 's ancestry such that  $K_{\text{contribs}} = r$ .
- For each  $C \in C_1 \dots C_n$ ,  $C_{\text{sig}}$  is a signature by the author  $\alpha$ .
- $C_1 \dots C_n$  have never been claimed before on  $P$ . Note that if two projects have a shared ancestry, it's possible to claim contributions once for *each* project.

## OUTPUTS

- $\alpha_{\text{bal}'} = \alpha_{\text{bal}} + v$ , where  $\alpha_{\text{bal}'}$  is the author's new account balance, and  $v = P_{\text{contract}}(t)$  where  $t$  is the claim-contribution transaction.